

Backtest Engineering: Designing a Correct Event-Driven Research Engine and Quantifying Performance Bias from Execution Simplifications: A Case Study on Ten Liquid US Equity ETFs

Srijan Gupta

Independent Researcher

srijangupta2013@gmail.com

February 2026

Abstract

A systematic researcher who reports a Sharpe ratio above 0.5 from a daily backtest and then deploys the strategy live is likely to be disappointed. The gap between paper performance and live performance is well documented in the practitioner literature, yet the specific magnitude of performance distortion caused by each individual execution simplification remains poorly characterised in the academic backtesting literature. This paper makes three contributions using a ten-instrument case study of liquid US equity sector ETFs (SPY, QQQ, IWM, DIA, XLF, XLK, XLE, XLV, XLY, XLP) over 2005–2025. First, we present a minimal, open-source event-driven backtesting engine whose design centres on strict causality enforcement, double-entry accounting invariants, and a deterministic four-stage event queue verified by a formal unit-test suite. Second, we introduce a six-level *execution realism ladder* spanning naïve next-open fills (M0), commission fees (M1), bid–ask spread (M2), volatility-scaled slippage (M3), a square-root market-impact proxy with participation-rate cap (M4), and delayed execution (M5), each implemented as a composable module with fully documented parameters. Third, using two canonical daily strategies, namely time-series momentum (TSMOM-60) and a z-score mean-reversion rule (MeanRev-z1)—the latter serving explicitly as a cost-sensitivity diagnostic—we measure how reported performance shifts across all six execution tiers. We report block-bootstrap 95% confidence intervals for each Sharpe estimate and conduct a sensitivity sweep over the impact-scaling parameter $k_{\text{imp}} \in \{0, 0.1, 0.25, 0.5, 0.75, 1.0, 1.5, 2.0\}$. Within this specific ten-ETF universe, the TSMOM-60 Sharpe ratio falls from 0.53 (M0, naïve) to 0.025 (M4, impact-constrained), a Sharpe inflation ratio of $\approx 21\times$; crucially, the naïve M0 Sharpe (0.53, CI $[-0.19, +0.94]$) is itself statistically indistinguishable from zero, as are all intermediate tiers (M1–M3), and the M4 CI $[-0.39, +0.47]$ confirms that realistic-execution performance is equally indistinguishable from zero. The MeanRev-z1 strategy, which earns a naïve Sharpe of 0.12 (CI $[-0.23, +0.52]$, itself spanning zero), collapses to -0.52 after adding fees alone, a sign flip that renders the naïve backtest directionally misleading and illustrates how per-tier cost attribution can diagnose non-viable strategies before capital commitment.

These results and the tooling used to produce them are offered as a reproducible reference for researchers seeking more credible empirical standards in systematic strategy evaluation.

Keywords: backtesting, execution realism, transaction costs, event-driven simulation, time-series momentum, mean reversion, Sharpe ratio, market impact, confidence intervals.

JEL: G12, G14, C63.

1 Introduction

Consider a practitioner who codes a daily momentum strategy, observes a Sharpe ratio near 0.5 on five years of ETF data, and allocates capital to it. Upon going live, the strategy earns a fraction of the backtested performance, or loses money altogether. This outcome is common. Practitioners attribute the shortfall to *execution slippage*, a broad term covering everything from the bid–ask spread on each trade to the price impact of moving a meaningful position through a thinly-traded order book.

Statistical sources of backtest inflation (multiple testing, overfitting, and data snooping) have received substantial academic attention. [Harvey et al. \[2016\]](#) argue that the t-statistic hurdle for new factors should be at least 3.0 to account for implicit testing. [Bailey et al. \[2014\]](#) formalise the Deflated Sharpe Ratio, which penalises reported performance by the number of strategies implicitly evaluated during selection. [White \[2000\]](#) and [Sullivan et al. \[1999\]](#) develop bootstrap tests for trading rule overfitting. What these contributions share is a focus on *statistical* inflation. The present paper addresses a complementary source: *execution-modeling inflation*, the distortion introduced when backtests simulate execution under assumptions that bear little resemblance to live trading.

Market microstructure theory has long established that execution is costly. [Kyle \[1985\]](#) demonstrates that informed trading generates endogenous price impact, with the resulting Kyle lambda parameterising linear (not square-root) price impact in a sequential rational-expectations auction model. The empirical square-root scaling of temporary execution impact is established in later work: [Almgren et al. \[2005\]](#) provide the primary empirical calibration using institutional trade records, while [Gatheral \[2010\]](#) establishes theoretical no-arbitrage conditions supporting concave impact functions. [Keim and Madhavan \[1997\]](#) document large institutional execution costs that vary systematically with order size and instrument liquidity. [Frazzini et al. \[2012\]](#) show that transaction costs materially erode the live performance of many well-known academic factor strategies. [Amihud \[2002\]](#) demonstrates that illiquidity is itself a priced risk factor.

Prior work has systematically characterised how transaction costs erode factor returns across large anomaly sets [[Novy-Marx and Velikov, 2016](#)]. These studies establish the broad empirical fact that many published anomalies are not exploitable after realistic costs. The present paper is complementary rather than redundant: it offers a transparent, modular, open-source engine with auditable cost components, per-tier attribution on a specific named universe, and a formal invariant test suite that verifies the correctness of the simulation itself—features that large-scale anomaly taxonomies, by construction, do not provide. The goal is to make credible cost-aware backtesting accessible without requiring production-scale infrastructure.

Despite this rich microstructure literature, many widely-used backtesting frameworks either ignore execution costs entirely or expose them through opaque interfaces that researchers apply

without examining the underlying assumptions. The present paper closes this gap through three contributions: an auditable event-driven engine with formal invariant tests, a transparent six-level cost ladder with a single authoritative parameter table, and a controlled quantification of execution-modeling inflation on a specific named ten-instrument universe (no extrapolation claimed).

The paper is organised as follows. Section 2 situates the work in the literature. Section 3 describes the engine. Section 4 formalises the six execution models. Section 5 specifies the experimental design. Section 6 presents results. Section 7 interprets findings. Section 8 acknowledges limitations. Section 9 concludes.

2 Related Work

2.1 Backtesting Bias and Multiple Testing

White [2000] and Sullivan et al. [1999] develop bootstrap-based tests showing that most technical trading rule excess returns vanish under proper multiple-comparison correction. Harvey et al. [2016] catalog over 300 published factors and recommend a t-statistic threshold of 3.0 against false discovery. Bailey et al. [2014] operationalise the cost of a strategy search via the Deflated Sharpe Ratio. These works address *statistical* inflation; the current paper addresses the complementary *economic* dimension from execution modeling.

2.2 Transaction Costs in Strategy Research

Keim and Madhavan [1997] provide early evidence that institutional execution costs are economically significant, with round-trip costs often exceeding 50–100 bps for mid-cap stocks. Lesmond et al. [1999] develop an implicit transaction cost estimator from the incidence of zero-return days, providing a parsimonious proxy for the cost levels that render many published strategies non-viable. Frazzini et al. [2012] analyse live institutional trade data and document that many academic factor strategies are barely profitable after realistic transaction costs. Amihud [2002] demonstrates that illiquidity predicts cross-sectional expected returns.

2.3 Price Impact Models

Kyle [1985] analyses a stylised sequential rational-expectations model in which a strategic informed trader causes *permanent* price impact proportional to signed order flow, a linear relationship, not square-root. The empirical square-root scaling of *temporary* execution impact, the relationship $\iota \propto \sqrt{Q/\text{ADV}}$, is established in later empirical work. Almgren et al. [2005] provide the primary calibration of this law from institutional trade records; Gatheral [2010] establishes theoretical

conditions under which concave impact is consistent with no-dynamic-arbitrage. [Almgren and Chriss \[2001\]](#) derive optimal liquidation trajectories under temporary-plus-permanent impact, providing the theoretical basis for participation-rate constraints.

2.4 Cost-Adjusted Strategy Performance

[Novy-Marx and Velikov \[2016\]](#) systematically study after-cost performance across a large set of anomalies, finding that strategies with monthly turnover above roughly 50% rarely survive realistic transaction costs. [Korajczyk and Sadka \[2004\]](#) test cross-sectional momentum strategies against intraday price-impact measures and estimate break-even fund sizes of \$5 billion or more for liquidity-weighted implementations. [Patton and Weller \[2020\]](#) extend Fama-MacBeth regressions to mutual fund data, delivering all-in cost estimates that eliminate on-paper momentum returns for typical funds. The present paper differs from these studies in three respects: it attributes cost impact at each individual friction tier (fees, spread, slippage, impact, delay) separately rather than as a lump sum; it provides a formally tested open-source engine as a reusable artefact; and it focuses on ETF strategies, where the liquidity regime differs materially from individual-stock anomaly portfolios.

2.5 Event-Driven Backtesting

Vectorised backtests iterate over a return matrix simultaneously, which can introduce look-ahead bias when the same bar's open and close prices are used in the same computation without explicit safeguards. Event-driven architectures [[Quantopian, Inc., 2016](#)] process market events sequentially, making execution timing explicit. [Lopez de Prado \[2018\]](#) advocates for event-driven design precisely because it forces the researcher to state timing assumptions. The engine described here follows this paradigm and adds a formal test suite to verify that the implemented invariants are actually maintained, a step that existing libraries do not uniformly provide.

Table 1 provides a structured comparison of this engine against four widely-used alternatives.

Table 1: Design comparison: this engine vs. established backtesting frameworks. “Composable cost tiers” means friction components can be independently activated. “Formal invariant tests” means the library ships unit tests verifying causality, accounting identity, and fill consistency. “Partial” indicates that the feature exists but requires user configuration and is not enforced by default.

Feature	This engine	Zipline	Backtrader	VectorBT	LEAN
Event-driven architecture	✓	✓	✓	✗	✓
Formal causality unit tests	✓	✗	✗	✗	✗
Accounting identity tests	✓	✗	✗	✗	✗
Composable cost tiers	✓	partial	partial	✗	partial
Partial-fill / ADV cap	✓	✗	✗	✗	partial
YAML reproducibility	✓	✗	✗	✗	partial
Block-bootstrap CI output	✓	✗	✗	✗	✗
Active maintenance (2025)	✓	✗	✓	✓	✓

Note: Feature comparisons are based on inspection of public repositories as of February 2026. VectorBT’s vectorised design precludes event-level causality tests by construction. LEAN (QuantConnect) has an extensive test infrastructure but does not expose backtesting-specific accounting-invariant verification in the sense defined here. The original Zipline library [[Quantopian, Inc., 2016](#)] is no longer actively maintained following Quantopian’s closure in 2020; a community fork (`zipline-reloaded`) exists but differs from the original. Readers should verify against the current version of each library at the respective repository.

2.6 Sharpe Ratio Inference under Autocorrelation

Sharpe [[1994](#)] defines the generalised Sharpe ratio used throughout this paper. Lo [[2002](#)] shows that standard Sharpe ratio standard errors are materially underestimated when returns are autocorrelated. The moving-blocks bootstrap of Künsch [[1989](#)], which is distinct from the later stationary bootstrap of Politis and Romano [[1994](#)] that uses random rather than fixed block lengths, provides a nonparametric approach to Sharpe confidence intervals that preserves serial dependence without imposing a distributional model.

3 Engine Design

3.1 Architecture

The engine contains four components connected through a typed event queue. A `DataHandler` is the sole interface for historical price data. A `Strategy` maps market observations to signals. A `Portfolio` converts signals into orders subject to cash and position constraints. An `ExecutionHandler` simulates the fill process according to the configured friction model. Within each trading day, these activate in the fixed order shown in Figure 1.




Figure 1: The four-stage FIFO event queue. Processing is strictly sequential within each trading day: no component at stage k may consume data that becomes available only at stage $k + 1$.

3.2 Causality Enforcement

The most common inadvertent look-ahead bias in daily backtests is computing a signal using closing price p_t^{close} then filling at the same day’s open p_t^{open} , which in OHLCV data is sequentially *earlier* than the close. The `DataHandler` prevents this by enforcing that any call to `get_history_asof(s, t)` returns price history indexed at all times $t' \leq t$ and strictly no data at $t' > t$, implemented via a per-symbol monotonically advancing timestamp pointer that prevents any symbol from accessing a future bar during the current day’s processing.

Data provenance caveat. This mechanism prevents *code-level* look-ahead but does not prevent the subtler issue of historically back-adjusted prices: if the data provider retroactively adjusts historical bars after corporate actions, the as-of slice will contain the adjusted history, which incorporates a future dividend event. The data validation described in Section 5 mitigates but does not entirely eliminate this concern.

3.3 Correctness Invariants and Test Suite

Four invariants are verified by automated unit tests:

- (i) **Accounting identity.** At every timestamp t :

$$E_t = C_t + \sum_s q_t^s \cdot p_t^{s,\text{close}}, \quad (1)$$

where E_t is equity, C_t is cash, q_t^s is shares held in symbol s , and $p_t^{s,\text{close}}$ is the closing price used for end-of-day mark-to-market. Fills execute at the next bar’s open; the mark-to-market therefore reflects the closing price on the signal day, before the fill is applied.

- (ii) **Strict causality.** Strategy computation during bar t may only access prices with index $t' \leq t$.
- (iii) **Fill quantity bound.** For every fill: $0 \leq q^{\text{filled}} \leq q^{\text{ordered}}$. Note that the impact-adjusted fill *price* can exceed the observed day’s high when large cost adjustments are applied; the invariant therefore constrains quantity, not price, relative to OHLC bounds.

- (iv) **Cost monotonicity.** In a controlled single-symbol long-only test on deterministic data, enabling any cost parameter cannot strictly increase total return.

3.4 Portfolio Mechanics

The portfolio uses a long-only, target-weight sizing rule. On a buy signal for symbol s at time t , the target position value is $w \cdot E_t$, clipped by available cash and a maximum-weight cap:

$$q_s^{\text{target}} = \left\lfloor \frac{\min(wE_t, w_{\max}E_t, C_t + q_t^s \cdot p_t^{s,\text{close}})}{p_t^{s,\text{close}}} \right\rfloor. \quad (2)$$

An order is generated only when $|\Delta q| \geq 1$. On a sell signal, the full open position is liquidated. In the experiments, $w = w_{\max} = 1.0$ (fully invested, no leverage).

4 Execution Realism Ladder

Table 2 is the single authoritative parameter specification for all six models. Each model is *cumulative*: Model k includes all parameters from Model $k - 1$ plus one additional friction component. **All results in this paper were generated from a single run using exactly the parameters in Table 2.** No table derives from a different run or parameter configuration.

Table 2: Authoritative parameter table for the six execution models. f_{fee} : commission per side (bps); h : half-spread per side (bps); k_{vol} : volatility-slippage scaling factor *in bps per unit of annualised decimal volatility* (so $k_{\text{vol}} = 10.0$ with $\sigma_{\text{ann}} = 0.20$ yields $s = 2.0$ bps); k_{imp} : impact scaling; ρ : participation-rate cap (fraction of 20-day ADV); δ : execution delay (trading days). $\delta = 1$ in all models except M5 means fills execute at the next bar's open ($t+1$ open); $\delta = 2$ in M5 means fills execute at the second next bar's open ($t+2$ open).

Tier	Model name	f_{fee}	h	k_{vol}	k_{imp}	ρ	δ
M0	Naïve	0	0	0	0	1.00	1
M1	+Fees	5	0	0	0	1.00	1
M2	+Spread	5	5	0	0	1.00	1
M3	+Slippage	5	5	10.0	0	1.00	1
M4	+Impact	5	5	10.0	0.50	0.05	1
M5	+Delay	5	5	10.0	0.50	0.05	2

M0: Naïve. $p^{\text{fill}} = p_{t+1}^{\text{open}}$, total cost = 0. This serves as the performance ceiling, not a realistic estimate.

M1: Commission Fees. fee = $(f_{\text{fee}}/10,000) \cdot p^{\text{fill}} \cdot q$. We use 5 bps as an *institutional cost proxy*; post-2019 retail ETF commissions at major US brokers are nominally zero, but clearing, market-data, and regulatory fees persist.

M2: Bid–Ask Spread. For a BUY order: $p^{\text{fill}} = p_{t+1}^{\text{open}} \cdot (1 + h/10,000)$. For a SELL order: $p^{\text{fill}} = p_{t+1}^{\text{open}} \cdot (1 - h/10,000)$. The label spread_10bps refers to the 10 bps *round-trip* (2 sides $\times h = 5$ bps). Quoted half-spreads for SPY are typically 1–2 bps; 5 bps is conservative for the full ten-ETF universe including thinner sector funds.

M3: Volatility-Scaled Slippage. $s_t = k_{\text{vol}} \cdot \sigma_{\text{ann},t}$ [bps], where σ_{ann} is the 20-day rolling annualised close-to-close volatility (decimal form, $n - 1$ degrees of freedom) and $k_{\text{vol}} = 10.0$. Fill price: $p^{\text{fill}} = p_{t+1}^{\text{open}} \cdot (1 \pm s_t/10,000)$, with + for BUY and – for SELL.

M4: Market Impact with Participation Cap. Using the empirical square-root law of [Almgren et al. \[2005\]](#):

$$\iota_t = k_{\text{imp}} \cdot \sqrt{V_{\text{trade}}/V_{\text{ADV}}} \quad [\text{fraction}], \quad (3)$$

where $V_{\text{trade}} = p_{t+1}^{\text{open}} \cdot q$ and V_{ADV} is the 20-day rolling average dollar volume. [Almgren et al. \[2005\]](#) report impact coefficients approximately in the range 0.3–0.8 for liquid equities; $k_{\text{imp}} = 0.50$ is near the lower-middle of empirically calibrated values.¹ The complete fill price, incorporating all cumulative costs, with explicit sign conventions for BUY and SELL directions, is:

$$p^{\text{fill}} = \begin{cases} p_{t+1}^{\text{open}} \cdot \left(1 + \frac{h}{10,000}\right) \cdot \left(1 + \frac{s_t + \iota_t^{\text{bps}}}{10,000}\right) & (\text{BUY}), \\ p_{t+1}^{\text{open}} \cdot \left(1 - \frac{h}{10,000}\right) \cdot \left(1 - \frac{s_t + \iota_t^{\text{bps}}}{10,000}\right) & (\text{SELL}), \end{cases} \quad (4)$$

where $\iota_t^{\text{bps}} = \iota_t \times 10,000$ and spread/impact are applied multiplicatively (compounded, not additive). The participation-rate cap limits fills to $\rho = 5\%$ of 20-day average daily share volume:

$$q^{\text{filled}} = \min(q^{\text{ordered}}, \lfloor \rho \cdot \text{ADV}_{\text{shares}} \rfloor). \quad (5)$$

We set $\rho = 5\%$ as a conservative institutional benchmark, consistent with sell-side execution desk guidelines for aggressive strategies. For SPY (daily volume > \$30B), this cap rarely binds for

¹The range 0.3–0.8 refers to the temporary-impact coefficient η in the parameterisation of [Almgren et al. \[2005\]](#), Table 2, which expresses impact as a fraction of bid–ask spread times $\sqrt{Q/V}$. The mapping to our k_{imp} (a dimensionless scaling on the square-root term) is approximate; researchers calibrating k_{imp} from TCA data should re-estimate using their own trade population.

typical retail-scale positions; for sector ETFs with daily volume of \$300–500M, it binds more frequently for large simulated positions. A sweep over $\rho \in \{0.05, 0.10, 0.20, 0.30\}$ would further quantify this effect and is recommended for future work. Section 6.5 reports a full sweep over k_{imp} .

M5: Delayed Execution. As in M4, but fills occur at $t + 2$ open ($\delta = 2$ trading days) rather than $t + 1$ open, modelling operational latency in a daily-signal workflow.

5 Experimental Setup

5.1 Data and Universe

Table 3: Instrument universe. Data sourced via Yahoo Finance adjusted close, incorporating dividends and splits, downloaded February 2026. *Revision note:* An earlier version of this project used Stooq as the data source; Yahoo Finance adjusted prices are used here. Adjusted prices incorporate the full corporate-action history as of the download date (a mild look-ahead bias standard in academic daily-bar research, acknowledged in Section 8).

Symbol	Name	Category	Inception
SPY	SPDR S&P 500 ETF	US broad equity	Jan 1993
QQQ	Invesco QQQ Trust	US tech/Nasdaq-100	Mar 1999
IWM	iShares Russell 2000 ETF	US small-cap	May 2000
DIA	SPDR Dow Jones Industrial ETF	US large-cap	Jan 1998
XLF	Financial Select SPDR	US financials	Dec 1998
XLK	Technology Select SPDR	US technology	Dec 1998
XLE	Energy Select SPDR	US energy	Dec 1998
XLV	Health Care Select SPDR	US health care	Dec 1998
XLY	Consumer Discretionary SPDR	US consumer disc.	Dec 1998
XLP	Consumer Staples SPDR	US consumer stat.	Dec 1998

All inception dates predate the sample start of January 2005; no proxy series is required. Data validation included: (i) checking for missing bars on US trading days (gaps forward-filled from prior close); (ii) excluding zero-volume days from ADV calculations; and (iii) inspecting single-bar return spikes exceeding 30% for data errors. No material quality issues were found for these ten instruments.

Universe selection bias. These ten instruments were selected to represent liquid, broad US equity market segments. All ten were continuously traded throughout the sample period and are not subject to individual-stock delisting risk. However, selecting *currently surviving and*

liquid instruments introduces universe-selection bias: the sample excludes ETFs launched and subsequently liquidated during this period. This is unlikely to reverse the directional findings, but results should not be interpreted as representative of the broader ETF universe.

5.2 Strategies

TSMOM-60. Following Moskowitz et al. [2012], the signal is:

$$r_{s,t}^{(60)} = p_{s,t}/p_{s,t-60} - 1. \quad (6)$$

Buy when $r_{s,t}^{(60)} > 0$; liquidate otherwise.

MeanRev-z1.

$$z_t = (r_t - \hat{\mu}_t)/\hat{\sigma}_t, \quad (7)$$

where $\hat{\mu}_t$ and $\hat{\sigma}_t$ are the sample mean and standard deviation ($n - 1$ d.f.) of the preceding 20 daily returns. Buy when $z_t < -1.0$; liquidate otherwise.

Diagnostic role. MeanRev-z1 is included explicitly as a cost-sensitivity diagnostic: because its naïve Sharpe is near zero and its turnover is approximately 180–200 round-trips per year, it is structurally predisposed to sign-flip under any meaningful per-trade cost. Its primary value is not as a viable strategy proposal but as a calibration instrument that reveals how quickly per-tier costs compound to produce directionally wrong naïve estimates.

Structural note on turnover. TSMOM-60 generates a signal change (buy to flat or flat to buy) when the 60-day return crosses zero, which occurs roughly 30–40 times per year in liquid equity ETFs, implying approximately 15–20 round-trips per year. MeanRev-z1 has a much higher turnover profile: the liquidation condition ($z_t \geq -1.0$) is satisfied on roughly 84% of trading days for standard normal returns. Empirically, MeanRev-z1 generates a new signal on approximately 90% of trading days, implying roughly 180–200 round-trips per year. This structural contrast explains why TSMOM remains viable through M3 (fees, spread, and slippage are applied only 15–20 times per year) while MeanRev is fatally sensitive to even the smallest per-trade cost.

5.3 Performance Metrics and Inference

$$\hat{S} = (\bar{r}/\hat{\sigma}) \cdot \sqrt{252} \quad (\text{annualised Sharpe ratio [Sharpe, 1994]}), \quad (8)$$

$$\text{CAGR} = (E_T/E_0)^{1/Y} - 1, \quad (9)$$

$$\text{MDD} = \min_t (E_t / \max_{s \leq t} E_s - 1). \quad (10)$$

Block-bootstrap 95% CIs for \hat{S} use the moving-blocks bootstrap of [Künsch \[1989\]](#) with block length $b = 10$ days, $N = 500$ resamples, and random seed 42 (ensuring exact reproducibility). We do *not* use the stationary bootstrap of [Politis and Romano \[1994\]](#), which draws random block lengths. As noted in Appendix B, $b = 10$ understates uncertainty for momentum strategies; reported CIs should be treated as lower bounds on true uncertainty. Wider blocks ($b = 20\text{--}30$) would only strengthen the conclusion that impact-constrained Sharpe estimates span zero.

5.4 Period Splits

Four evaluation windows: full sample (2005–2025), pre-/during-GFC (2005–2012), mid-cycle (2013–2019), and COVID-era and recovery (2020–2025). Note that the 2020–2025 window includes the sharp market decline of March 2020; it is not a post-crisis-only sample.

5.5 Buy-and-Hold Benchmark

An equally-weighted buy-and-hold portfolio over all ten instruments, rebalanced annually, is evaluated under M0 and M4. Its M0 Sharpe is approximately 0.48 (CAGR $\approx 10\%$), providing a passive baseline.

6 Results

6.1 Master Results Table

Table 4 is the single canonical results table for this paper. Every Sharpe value, CI, CAGR, and MDD cited in the text is sourced exclusively from this table, which derives from one reproducible run with the parameters in Table 2.

Table 4: Annualised Sharpe ratio, 95% block-bootstrap CI, CAGR, and maximum drawdown for all strategy–execution-model combinations (full sample, 2005–2025; single canonical run; bootstrap seed = 42). \dagger : CI spans zero; Sharpe statistically indistinguishable from zero at 95% level. \ddagger : upper CI bound marginally negative (-0.003), indicating the sign flip is robustly negative at the block length used; note however that this upper bound is sensitive to block length—wider blocks ($b = 20\text{--}30$) may shift the upper bound to slightly positive, so this result should be interpreted with caution. *Note:* All TSMOM rows M0–M5 have CIs spanning zero and therefore carry \dagger ; this notation is applied consistently to every row whose 2.5% CI bound is negative. Block length $b = 10$ understates uncertainty for momentum strategies; all CIs should be treated as *lower bounds* on true uncertainty. MDD = -0.998 denotes near-total (99.8%) capital loss. B&H = equally-weighted buy-and-hold benchmark.

Strategy	Model	Sharpe	CI [2.5%, 97.5%]	CAGR	MDD
B&H	M0: Naïve	0.480	[+0.12, +0.84]	+0.100	-0.510
TSMOM-60	M0: Naïve	0.530	[−0.19, +0.94] \dagger	+0.079	-0.418
TSMOM-60	M1: +Fees	0.479	[−0.22, +0.89] \dagger	+0.069	-0.429
TSMOM-60	M2: +Spread	0.435	[−0.22, +0.81] \dagger	+0.061	-0.438
TSMOM-60	M3: +Slippage	0.417	[−0.22, +0.79] \dagger	+0.058	-0.446
TSMOM-60	M4: +Impact	0.025	[−0.39, +0.47] \dagger	-0.011	-0.704
TSMOM-60	M5: +Delay	-0.007	[−0.45, +0.38] \dagger	-0.017	-0.727
MeanRev-z1	M0: Naïve	0.129	[−0.23, +0.52] \dagger	+0.008	-0.303
MeanRev-z1	M1: +Fees	-0.521	[−0.93, −0.003] \ddagger	-0.072	-0.799
MeanRev-z1	M2: +Spread	-1.052	[−1.63, −0.37]	-0.133	-0.952
MeanRev-z1	M3: +Slippage	-1.227	[−1.78, −0.49]	-0.152	-0.970
MeanRev-z1	M4: +Impact	-2.246	[−2.85, −1.40]	-0.267	-0.999
MeanRev-z1	M5: +Delay	-1.697	[−2.39, −0.95]	-0.254	-0.998

6.2 TSMOM-60: Monotonic Degradation

Time-series momentum degrades monotonically as execution assumptions become more realistic (Table 4). Under M0, TSMOM-60 reports Sharpe 0.530 and CAGR 7.9%. Notably, this falls below the equally-weighted buy-and-hold benchmark’s CAGR of approximately 10%, meaning the active strategy underperforms passive investment even before realistic costs are applied. Practitioners should interpret a naïve Sharpe of 0.530 as evidence of a positive risk-adjusted gross return, not alpha.

Fees (M1) reduce Sharpe to 0.479; adding spread (M2) yields 0.435; incorporating volatility slippage (M3) yields 0.417. Importantly, all three tiers also have CIs spanning zero (M1: $[-0.22, +0.89]$; M2: $[-0.22, +0.81]$; M3: $[-0.22, +0.79]$), as does M0 itself ($[-0.19, +0.94]$).

TSMOM-60’s Sharpe ratio is statistically indistinguishable from zero at the 95% confidence level under every execution model, including naïve assumptions. The wide CIs reflect the high variance of daily momentum returns over a 20-year sample and the conservative $b = 10$ block length, which understates uncertainty for autocorrelated return series.

The impact proxy with participation cap (M4) produces a sharp discontinuity: Sharpe falls from 0.417 to 0.025, CAGR turns negative (-1.1%), and maximum drawdown deteriorates to -70.4% . The Sharpe inflation ratio is:

$$\text{IR} = \hat{S}_{M0}/\hat{S}_{M4} = 0.530/0.025 \approx 21.2 \times . \quad (11)$$

This ratio is specific to the parameters $k_{\text{imp}} = 0.50$ and $\rho = 0.05$; the sensitivity sweep in Section 6.5 and the ρ robustness check recommended in Section 8 provide the necessary context for interpreting its magnitude. Critically, the M4 Sharpe of 0.025 has a 95% bootstrap CI of $[-0.39, +0.47]$, which spans zero by a wide margin. **Under impact-constrained execution, TSMOM-60’s performance is statistically indistinguishable from zero at the 95% confidence level.** The naïve backtest suggests a viable strategy; the realistic backtest reveals that the positive point estimate is statistical noise within a very wide uncertainty band.

The M4 discontinuity arises from three interacting mechanisms. First, the square-root impact law is superlinear: doubling the traded fraction of ADV more than doubles the cost. Second, the 5% ADV participation cap converts intended position changes into multi-day ladders of partial fills, each independently incurring spread, slippage, and impact. Third, during the fill accumulation period, the strategy holds an undersized position that generates signal costs without full position payoff, compounding the drag. Together, these produce a qualitatively different cost regime compared to M0–M3.

6.3 MeanRev-z1: Diagnostic of Cost Sensitivity

Under M0, MeanRev-z1 Sharpe is 0.129 with CI $[-0.23, +0.52]$. **The naïve Sharpe itself spans zero: even without transaction costs, MeanRev-z1’s positive Sharpe is statistically indistinguishable from zero.** The 20-day z-score rule on daily returns of liquid broad-market ETFs confronts strong arbitrage forces [Jegadeesh, 1990, Gatev et al., 2006].

Adding fees alone (M1) drives Sharpe from 0.129 to -0.521 , a sign reversal. This result illustrates the diagnostic value of the tier-by-tier ladder: a cost that any live trader immediately prices in causes the expected Sharpe to flip sign. The structural explanation is the strategy’s approximately 180–200 round-trips per year: each application of a 5-bps fee compounds to eliminate far more than the strategy’s negligible per-trade gross edge. Researchers can use this pattern—a sign flip at the very first cost tier—as a strong signal that the naïve Sharpe is providing qualitatively wrong

information, warranting explicit cost quantification before any capital commitment is considered.

Adding spread (M2) produces -1.052 ; slippage (M3) yields -1.227 ; impact (M4) drives Sharpe to -2.246 with maximum drawdown approaching 99.8% of initial capital. A MDD of -0.998 represents near-total capital loss, an outcome requiring no further statistical analysis to characterise as non-viable.

M4-to-M5 non-monotonicity. MeanRev Sharpe improves from -2.246 (M4) to -1.697 (M5) when a two-day execution delay is added. This is counterintuitive but structurally explicable: the forced delay effectively reduces strategy turnover. When execution is delayed by one day, signals that would trigger an immediate trade-and-liquidation cycle may expire or reverse before being filled, bypassing some round-trip cost applications. For a high-turnover mean-reversion strategy, the turnover reduction from the delay dominates the cost of executing at a less-timely price. This mechanism is specific to long-only, liquidation-on-signal strategies with high daily signal frequency; it would not arise in persistent-position or long-short strategies.

6.4 Inflation Ratio Summary

Table 5: Sharpe inflation ratios sourced from Table 4 (single canonical run). $IR = \hat{S}_{M0}/\hat{S}_{Mk}$, defined only when both Sharpes share positive sign. For MeanRev, the sign flip occurs at M1; actual Mk Sharpe values are shown in brackets.

Strategy	Model	IR	\hat{S}_{Mk}
TSMOM-60	M1: +Fees	$1.1\times$	+0.479
TSMOM-60	M2: +Spread	$1.2\times$	+0.435
TSMOM-60	M3: +Slippage	$1.3\times$	+0.417
TSMOM-60	M4: +Impact	$21.2\times$	+0.025
TSMOM-60	M5: +Delay	sign flip	[−0.007]
MeanRev-z1	M1: +Fees	sign flip	[−0.521]
MeanRev-z1	M2: +Spread	sign flip	[−1.052]
MeanRev-z1	M3: +Slippage	sign flip	[−1.227]
MeanRev-z1	M4: +Impact	sign flip	[−2.246]

6.5 Impact Parameter Sensitivity

Figure 2 plots annualised Sharpe ratio against $k_{imp} \in \{0, 0.1, 0.25, 0.5, 0.75, 1.0, 1.5, 2.0\}$ for both strategies, with all other M4 parameters held constant. TSMOM degrades monotonically from its M3 Sharpe (≈ 0.42 at $k_{imp} = 0$) toward negative values as impact scales up. Even at $k_{imp} = 0.10$ TSMOM Sharpe is materially below the naïve value, confirming that the participation cap alone

contributes meaningfully at lower impact magnitudes. MeanRev is negative at all k_{imp} values because its sign flip occurs from fees alone, independently of impact.




Figure 2: Annualised Sharpe ratio vs. impact scaling parameter k_{imp} (M4 model, $\rho = 0.05$ held constant). TSMOM degrades monotonically from its slippage-only baseline toward negative Sharpe. MeanRev is negative at all impact levels because its sign flip occurs from fees alone. The baseline $k_{\text{imp}} = 0.50$ used in main results is marked with a vertical dotted line.

6.6 Period-Split Results

Table 6: Period-split Sharpe ratios for M0 (naïve), M1 (+fees), and M4 (+impact). The M4 Sharpe is lower than M0 in every sub-period for TSMOM, confirming that the degradation is not regime-specific. *Bootstrap CIs are not reported for sub-periods*: the shorter windows (6–8 years) produce materially wider uncertainty than the full-sample CIs in Table 4; all sub-period point estimates should be treated with correspondingly greater caution.

Strategy	Period	M0 Sharpe	M1 Sharpe	M4 Sharpe
TSMOM-60	Full (2005–2025)	+0.530	+0.479	+0.025
TSMOM-60	2005–2012	+0.043	+0.011	-0.132
TSMOM-60	2013–2019	+0.612	+0.558	+0.108
TSMOM-60	2020–2025 (COVID-era)	+0.959	+0.894	+0.317
MeanRev-z1	Full (2005–2025)	+0.129	-0.521	-2.246
MeanRev-z1	2005–2012	+0.149	-0.524	-2.441
MeanRev-z1	2013–2019	+0.087	-0.493	-2.197
MeanRev-z1	2020–2025 (COVID-era)	+0.143	-0.541	-2.184

Sub-period interpretation. The 2005–2012 TSMOM M0 Sharpe of 0.043 is low because the 2007–2009 global financial crisis created a sharp, correlated equity selloff in which all ten US equity instruments declined together. A long-only momentum rule in a single-asset-class universe has no cross-asset hedge and is directionally wrong throughout a sharp drawdown. This contrasts with multi-asset time-series momentum [Moskowitz et al., 2012], where crisis-period returns benefit from diversification that an equity-only universe cannot provide. The 2020–2025 TSMOM M0 Sharpe of 0.959 reflects unusually strong trend conditions: the rapid recovery of 2020–2021, the broad equity selloff of 2022 driven by Federal Reserve rate tightening, and the subsequent multi-year recovery. Trending markets with persistence beyond 60 trading days are the natural habitat of a 60-day lookback momentum rule. Across all four sub-periods, the qualitative finding holds: adding M4 execution materially reduces Sharpe relative to M0, with the magnitude of degradation largest in 2005–2012 where M0 Sharpe itself is near zero.

6.7 Visual Summaries

Figures 3 through 5 provide visual summaries fully consistent with Table 4.

Sharpe Ratio by Execution Model




Figure 3: Annualised Sharpe ratio by execution model, shown separately for each strategy. Error bars are 95% block-bootstrap CIs ($N = 500, b = 10$). TSMOM degrades monotonically; MeanRev transitions to negative at M1. Blue bars indicate non-negative Sharpe; red bars indicate negative Sharpe.




Figure 4: Sharpe inflation ratio for TSMOM-60, defined as $\hat{S}_{M0}/\hat{S}_{Mk}$ only when both Sharps are positive. M5 is marked as sign-flip (TSMOM M5 Sharpe = -0.007). MeanRev inflation ratios are omitted because a sign flip occurs at M1. The $21.2 \times$ M4 ratio ($= 0.530/0.025$, consistent with Table 5) reflects compounding of square-root impact and the 5% ADV participation cap.




Figure 5: Sharpe ratio heatmap. Columns are ordered M0 through M5; cells are annotated with numeric Sharpe values from Table 4. Diverging colour scale (green = positive, red = negative) makes the MeanRev sign flip at M1 immediately visible.

7 Discussion

7.1 The Passive Benchmark Context

The naïve TSMOM strategy reports CAGR 7.9%, below the equally-weighted buy-and-hold benchmark’s approximately 10% over the same period. Under realistic execution (M4), TSMOM CAGR turns negative (-1.1%). Even before asking whether the active strategy beats passive, the researcher must contend with the fact that the naïve Sharpe of 0.530 overstates impact-constrained performance by a factor of 21.

Notably, the equally-weighted buy-and-hold benchmark is the *only* entry in Table 4 with a confidence interval entirely above zero: CI $[+0.12, +0.84]$. Passive investment over this universe and period is statistically distinguishable from zero performance; TSMOM-60 under any execution assumption is not. This makes the passive benchmark, rather than a zero-Sharpe null, the appropriate comparison point for evaluating whether any active strategy adds value.

7.2 Why Impact Creates a Discontinuity

The performance drop between M3 and M4 results from three interacting mechanisms: (i) super-linear square-root impact cost; (ii) the 5% ADV cap converting position changes into multi-day partial fill sequences; and (iii) partial fill accumulation, wherein the strategy pays costs proportional to the intended position size but earns returns proportional to the smaller filled position. Together, these create a qualitatively different cost regime that cannot be approximated by scaling the linear M1–M3 costs.

7.3 Per-Tier Attribution as a Research Diagnostic

The sign flip in MeanRev at M1 illustrates the primary practical value of the execution realism ladder: by attributing performance separately at each tier, a researcher can immediately identify which cost channel is fatal to a strategy and why. When adding a 5-bps institutional fee causes the expected Sharpe to flip sign, the naïve backtest is providing qualitatively wrong information. The root cause is structural: approximately 180–200 round-trips per year at 5 bps each removes far more from returns than the strategy’s negligible gross edge per trade provides. More broadly, any strategy whose Sharpe point estimate is near zero under naïve assumptions should be tested against the full cost ladder before any further development effort—not because such strategies are necessarily unviable, but because their cost sensitivity cannot be inferred from the naïve Sharpe alone. The ladder makes this test straightforward and reproducible.

7.4 Statistical Interpretation

The bootstrap CIs in Table 4 carry an important message beyond any individual strategy result: *every* Sharpe point estimate in this study is either statistically indistinguishable from zero or robustly negative. TSMOM M4 CI [−0.39, +0.47] spans zero; MeanRev M0 CI [−0.23, +0.52] spans zero; and even the naïve TSMOM CIs (M0–M3) all have negative lower bounds. **A Sharpe point estimate reported without a confidence interval is incomplete information.** The 20-year sample at daily frequency, with the autocorrelation structure of these returns and a conservative block length, does not provide sufficient statistical power to confidently distinguish any active strategy from chance. Practitioners should report CIs alongside every point estimate; the open-source bootstrap implementation in this paper makes this straightforward.

8 Limitations

Daily-bar granularity. Friction models are stylised proxies that cannot represent intraday dynamics such as queue priority, limit order book depth, or the intraday profile of market impact. Tick-level simulation is required for more precise cost estimation.

Impact model calibration. The $21.2 \times$ inflation ratio for TSMOM is conditional on $k_{\text{imp}} = 0.50$ and $\rho = 0.05$. The sensitivity sweep (Figure 2) shows material variation across the sweep range. This figure should be interpreted as a specific case study, not a universal law.

Participation cap. We set $\rho = 5\%$ as a conservative institutional benchmark. A sweep over $\rho \in \{0.05, 0.10, 0.20, 0.30\}$ would further characterise the cap’s contribution and is recommended

for future work. The direction is predictable: a larger ρ reduces the participation constraint, reducing the multi-day fill laddering that drives the M4 discontinuity, and would therefore lower the $21.2 \times$ inflation ratio.

Universe, period, and strategy parameterisation. Results derive from ten US equity ETFs over 2005–2025. Lookback sensitivity for TSMOM (20, 40, 60, 120-day) and z-score threshold sensitivity for MeanRev ($z \in \{0.5, 1.0, 1.5, 2.0\}$) are important ablations for future work.

Long-only constraint. Full liquidation on every sell signal creates artificially high turnover for MeanRev, likely overstating its cost sensitivity relative to a persistent-position or long-short implementation. This is a deliberate design choice for the diagnostic use case, but it limits the generalisability of the MeanRev cost estimates.

Data adjustment. Total-return adjusted prices incorporate the full corporate-action history as of the download date, which is not available in real time. This constitutes a form of look-ahead bias that is standard in academic daily-bar research but should be acknowledged.

Autocorrelation-adjusted Sharpe. Lo [2002] shows that the uncorrected annualised Sharpe (scaled by $\sqrt{252}$) overstates the information ratio when returns exhibit positive autocorrelation. The block-bootstrap CIs partially address this, but point estimates themselves are unadjusted. An autocorrelation correction for TSMOM would widen CIs and lower the effective information ratio, only strengthening the conclusion that M4 performance spans zero.

9 Conclusion

The opening thought experiment (a backtested Sharpe of 0.5 that evaporates in live trading) is not a rare failure mode. It follows predictably from the combination of near-zero edge per trade, high turnover, and execution assumptions that ignore the primary cost channels of institutional trading. This paper demonstrates, on a specific ten-instrument case study, that the inflation ratio between naïve and impact-constrained execution can exceed $21 \times$ for a strategy with genuine momentum signal (TSMOM-60), and that even this inflated baseline Sharpe becomes statistically indistinguishable from zero after proper confidence interval construction. For a cost-sensitivity diagnostic strategy (MeanRev-z1), the naïve Sharpe itself spans zero, and the sign of expected performance reverses at the first cost tier.

These findings depend on the specific parameters and universe used and should not be generalised without further empirical validation. They do, however, support clear methodological

recommendations. First: a Sharpe point estimate reported without a confidence interval is incomplete information. In this study, every active strategy Sharpe is statistically indistinguishable from zero under block-bootstrap inference, a finding that would be invisible without CIs. Second: daily backtests should routinely report results at multiple execution tiers, not only under naïve assumptions. Third: the passive buy-and-hold benchmark should always be reported alongside active strategy performance; it is the only Sharpe in this study with a CI entirely above zero. The open-source engine and experimental harness provided here are intended to make these practices accessible without requiring production-scale simulation infrastructure.

Reproducibility. All results can be reproduced from the project repository by running `make data && make run && make paper` after installing dependencies via `pip install -r requirements.txt`. Key dependency versions: Python 3.12; pandas 2.2; numpy 1.26; scipy 1.13; yfinance 0.2; matplotlib 3.8 (see `requirements.txt` for the full pinned specification). The block-bootstrap uses random seed 42 throughout; all other computations are deterministic. Yahoo Finance adjusted prices change retroactively with each download; results were generated from a data snapshot downloaded in February 2026. A frozen copy of this data snapshot is provided as a supplementary archive (`data_snapshot_2026-02.tar.gz`) to ensure exact reproducibility independent of future Yahoo Finance adjustments. Repository: <https://github.com/srijan-gupta/backtest-engine-paper> (to be made public upon acceptance).

References

- Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *Journal of Risk*, 3(2):5–39, 2001.
- Robert Almgren, Chee Thum, Emmanuel Hauptmann, and Hong Li. Direct estimation of equity market impact. *Risk*, 18(7):58–62, 2005.
- Yakov Amihud. Illiquidity and stock returns: Cross-section and time-series effects. *Journal of Financial Markets*, 5(1):31–56, 2002.
- David H. Bailey, Jonathan Borwein, Marcos Lopez de Prado, and Qiji Jim Zhu. The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107, 2014.
- Andrea Frazzini, Ronen Israel, and Tobias J. Moskowitz. Trading costs of asset pricing anomalies. Technical report, Chicago Booth Research Paper No. 14-05; AQR Capital Management, 2012.

Working paper; circulated since 2012, updated 2020. Available at SSRN: <https://ssrn.com/abstract=2294498>.

Evan Gatev, William N. Goetzmann, and K. Geert Rouwenhorst. Pairs trading: Performance of a relative-value arbitrage rule. *Review of Financial Studies*, 19(3):797–827, 2006.

Jim Gatheral. No-dynamic-arbitrage and market impact. *Quantitative Finance*, 10(7):749–759, 2010.

Campbell R. Harvey, Yan Liu, and Heqing Zhu. . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1):5–68, 2016.

Narasimhan Jegadeesh. Evidence of predictable behavior of security returns. *Journal of Finance*, 45(3):881–898, 1990.

Donald B. Keim and Ananth Madhavan. Transactions costs and investment style: An inter-exchange analysis of institutional equity trades. *Journal of Financial Economics*, 46(3):265–292, 1997.

Robert A. Korajczyk and Ronnie Sadka. Are momentum profits robust to trading costs? *Journal of Finance*, 59(3):1039–1082, 2004.

Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *Annals of Statistics*, 17(3):1217–1241, 1989.

Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.

David A. Lesmond, Joseph P. Ogden, and Charles A. Trzcinka. A new estimate of transaction costs. *Review of Financial Studies*, 12(5):1113–1141, 1999. doi: 10.1093/rfs/12.5.1113.

Andrew W. Lo. The statistics of Sharpe ratios. *Financial Analysts Journal*, 58(4):36–52, 2002.

Marcos Lopez de Prado. *Advances in Financial Machine Learning*. Wiley, Hoboken, NJ, 2018.

Tobias J. Moskowitz, Yao Hua Ooi, and Lasse Heje Pedersen. Time series momentum. *Journal of Financial Economics*, 104(2):228–250, 2012.

Robert Novy-Marx and Mihail Velikov. A taxonomy of anomalies and their trading costs. *Review of Financial Studies*, 29(1):104–147, 2016. doi: 10.1093/rfs/hhv063.

Andrew J. Patton and Brian M. Weller. What you see is not what you get: The costs of trading market anomalies. *Journal of Financial Economics*, 137(2):515–549, 2020. doi: 10.1016/j.jfineco.2020.02.012.

Dimitris N. Politis and Joseph P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994.

Quantopian, Inc. Zipline: A Pythonic algorithmic trading library. <https://github.com/quantopian/zipline>, 2016. Open-source software. Accessed February 2026. Note: Quantopian ceased operations in 2020; a community fork (`zipline-reloaded`) is available at <https://github.com/stefan-jansen/zipline-reloaded>.

William F. Sharpe. The Sharpe ratio. *Journal of Portfolio Management*, 21(1):49–58, 1994.

Ryan Sullivan, Allan Timmermann, and Halbert White. Data-snooping, technical trading rule performance, and the bootstrap. *Journal of Finance*, 54(5):1647–1691, 1999.

Halbert White. A reality check for data snooping. *Econometrica*, 68(5):1097–1126, 2000.

A Engine Loop Pseudocode

Algorithm 1 Event-Driven Backtest Loop

```

1: Input:  $\mathcal{D}$  (price data),  $\mathcal{S}$  (strategy),  $\mathcal{P}$  (portfolio),  $\mathcal{E}$  (execution)
2: for each trading day  $t$  do
3:    $\mathcal{P}.\text{update\_prices}(t, \mathcal{D})$ 
4:    $orders \leftarrow []$ 
5:   for each symbol  $s$  do
6:      $e_M \leftarrow \text{MarketEvent}(t, s)$ 
7:      $e_\sigma \leftarrow \mathcal{S}.\text{on\_market}(e_M, \mathcal{D}.\text{asof}(s, t))$             $\triangleright$  Stage 2: uses data  $\leq t$  only
8:     if  $e_\sigma \neq \emptyset$  then
9:        $e_O \leftarrow \mathcal{P}.\text{on\_signal}(e_\sigma)$ 
10:      if  $e_O \neq \emptyset$  then
11:         $orders.append(e_O)$ 
12:      end if
13:    end if
14:   end for
15:   (All orders batched before any fills; prevents fill for symbol A affecting available cash for symbol B.)
16:   for each  $e_O$  in  $orders$  do
17:      $e_F \leftarrow \mathcal{E}.\text{execute}(e_O, \mathcal{D})$                           $\triangleright$  Fill at  $t + \delta$  open
18:     if  $e_F \neq \emptyset$  then
19:        $\mathcal{P}.\text{on\_fill}(e_F)$ 
20:     end if
21:   end for
22:    $\mathcal{P}.\text{mark\_to\_market}(t, \mathcal{D})$ 
23: end for

```

B Moving-Block Bootstrap Procedure

Let $\{r_1, \dots, r_T\}$ be the daily return series. The 95% CI for \hat{S} is constructed by (i) drawing $k = \lceil T/b \rceil$ start indices s_1, \dots, s_k i.i.d. Uniform $\{1, \dots, T-b+1\}$ with replacement; (ii) forming pseudo-sample $r^* = [r_{s_1}, \dots, r_{s_1+b-1}, \dots]$ truncated to length T ; (iii) computing \hat{S}^* ; (iv) repeating $N = 500$ times with random seed 42 and reporting $[\hat{S}_{(0.025)}^*, \hat{S}_{(0.975)}^*]$. Block length $b = 10$ preserves short-range autocorrelation [Künsch, 1989]. For momentum strategies with positive return autocorrelation [Lo, 2002], block lengths of $b = 20$ or $b = 30$ would produce wider and

more conservative intervals. Because $b = 10$ understates uncertainty for TSMOM-60, the reported CIs should be treated as *lower bounds*: wider blocks would only strengthen the conclusion that impact-constrained Sharpe estimates are not statistically distinguishable from zero.